



# Data engineering on Azure Databricks - workshop

Anagha Khanolkar  
Azure Cloud Solution Architects – Data and AI

# About the lab

# Learn how to...

- Provision Azure resources – blob storage, Azure SQL DB, Databricks
- Learn how to copy a public dataset into Azure
- Learn how to build a data engineering pipeline in Azure Databricks
- Learn how to automate a report jobset that integrates reports generated in Databricks into an RDBMS

# The datasets

## NYC taxi data



Transactional data



Yellow taxi trips (675 million) | 2009 - 2017

Green taxi trips (59 million) | 2013 - 2017

CSV

Schema varies between taxi types

Schema varies across years

~30GB raw



Reference data



Trip months

Payment type

Rate code

Taxi zone

Trip type

Vendor

# The Azure services

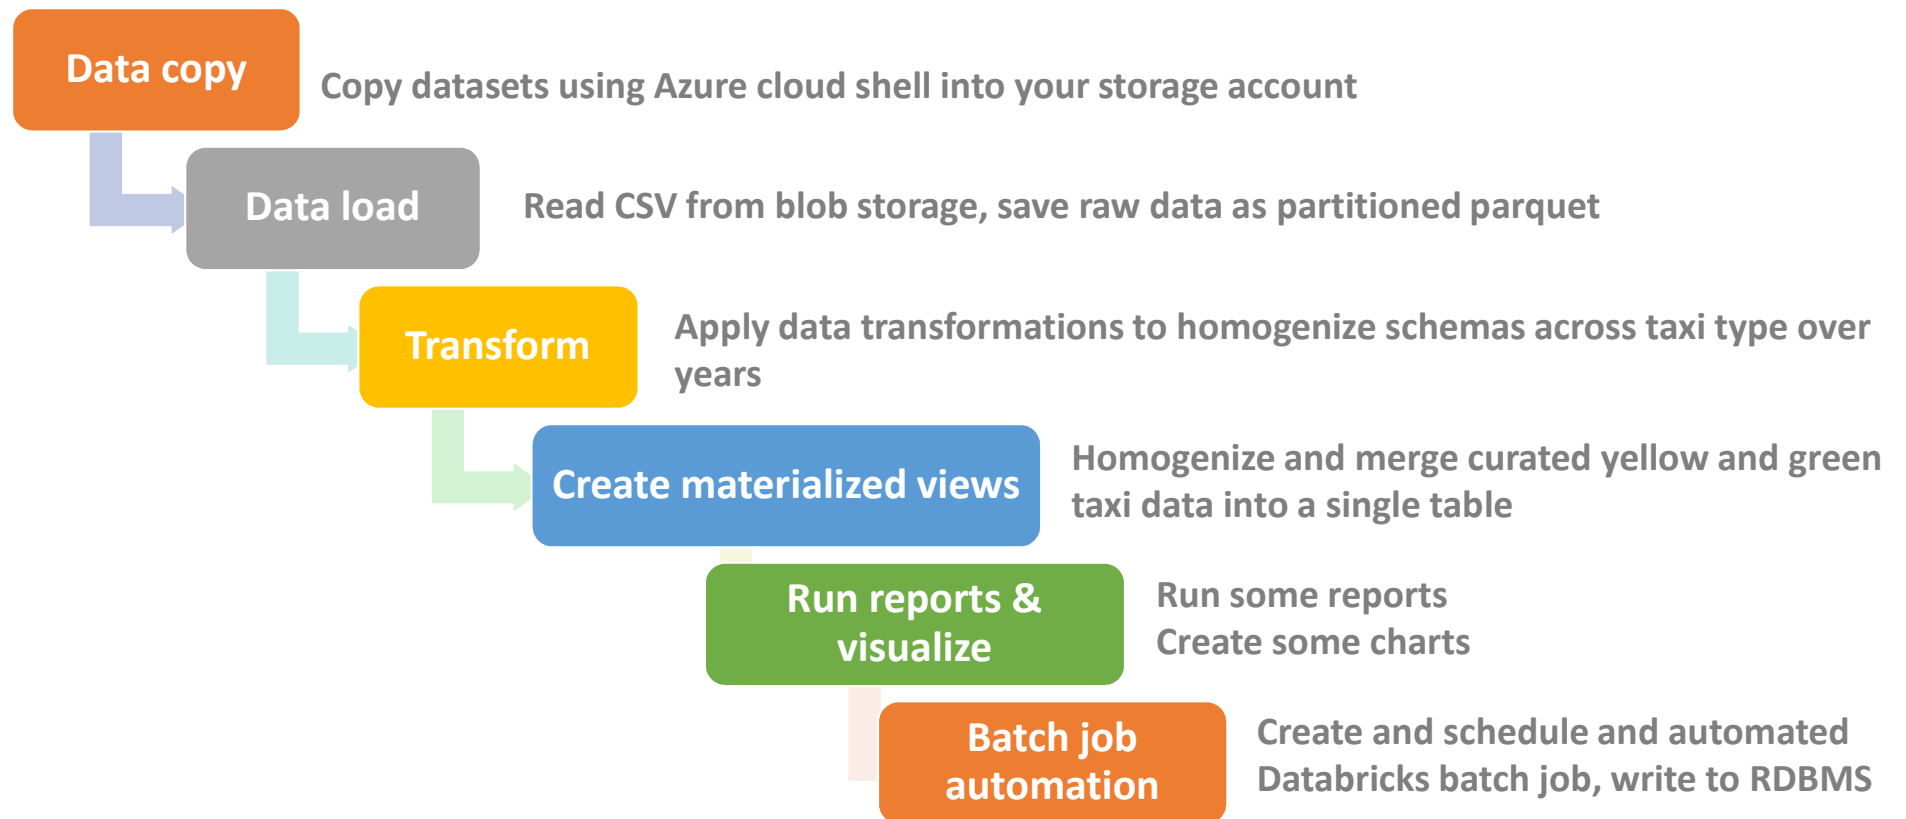


## Utilities:

Azure Cloud Shell – for provisioning storage and copying data

Azure SQL Database Query Explorer – cloud IDE for Azure SQL DB

# The lab modules



# **Data copy**

## **With Azure Cloud Shell - Bash**

# Data load TODOs

- ☐ Launch cloud shell
- ☐ Create a resource group in US East 2
- ☐ Create a storage account
- ☐ Create blob storage containers
- ☐ Copy data for the workshop (details in doc)
- ☐ Upload a subset of the workshop data into staging directory



## Instructions:

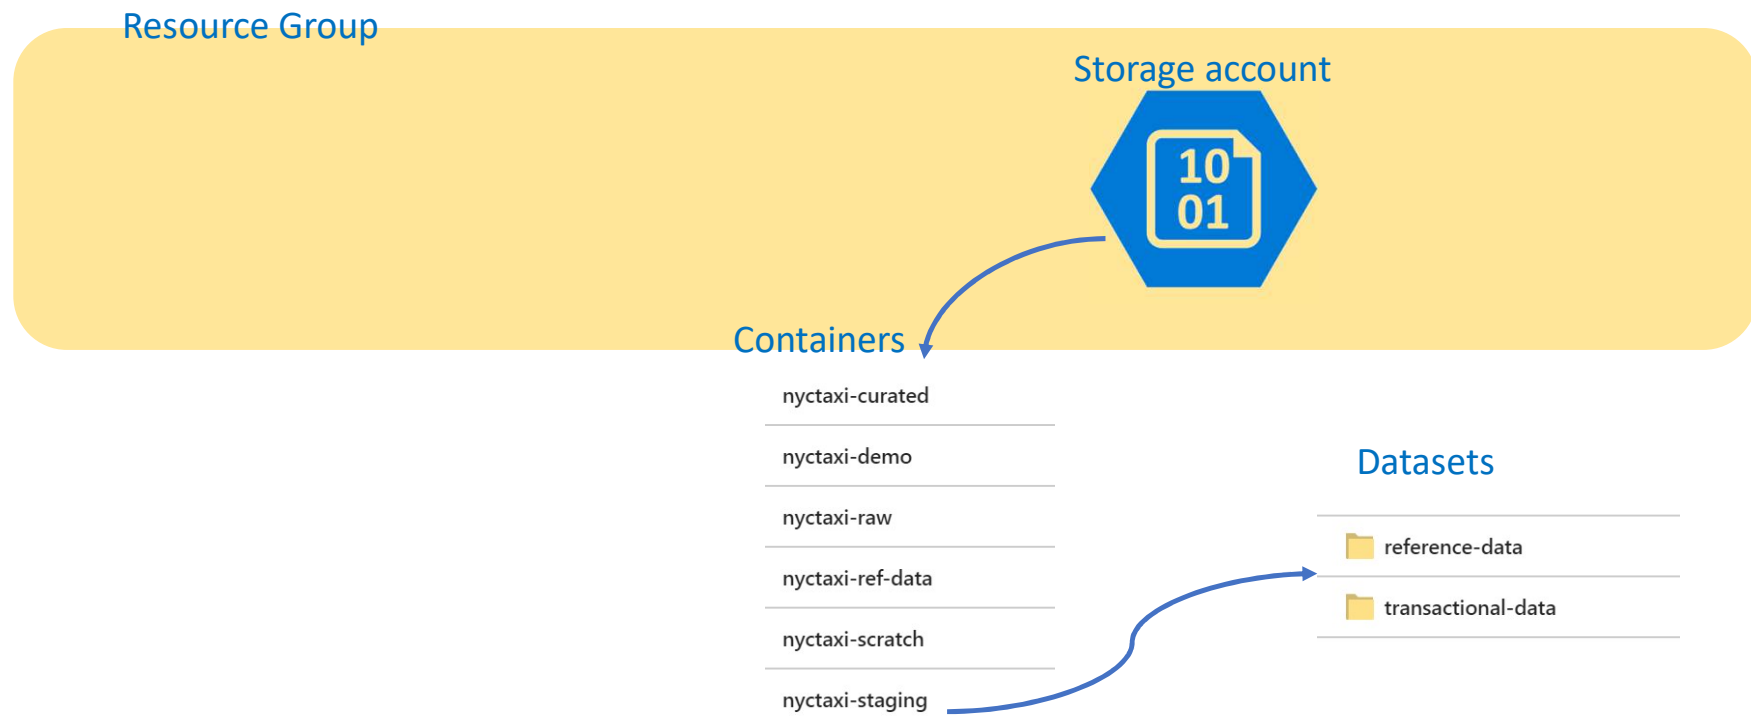
<https://aka.ms/DEW-Setup>

Copying the entire dataset (30 GB) took 2 hours USE2-USE2  
For the lab- we will use a subset



# By now...

You should have the following set up in your subscription:



# **Azure services**

## **Provisioning, configuration**

# Provisioning TODOs

- ☐ Provision Azure SQL Database in the resource group (100 DTU)
- ☐ Create tables in the Azure SQL Database – for batch job history & reports
- ☐ Provision Azure Databricks workspace in the resource group
- ☐ Launch workspace

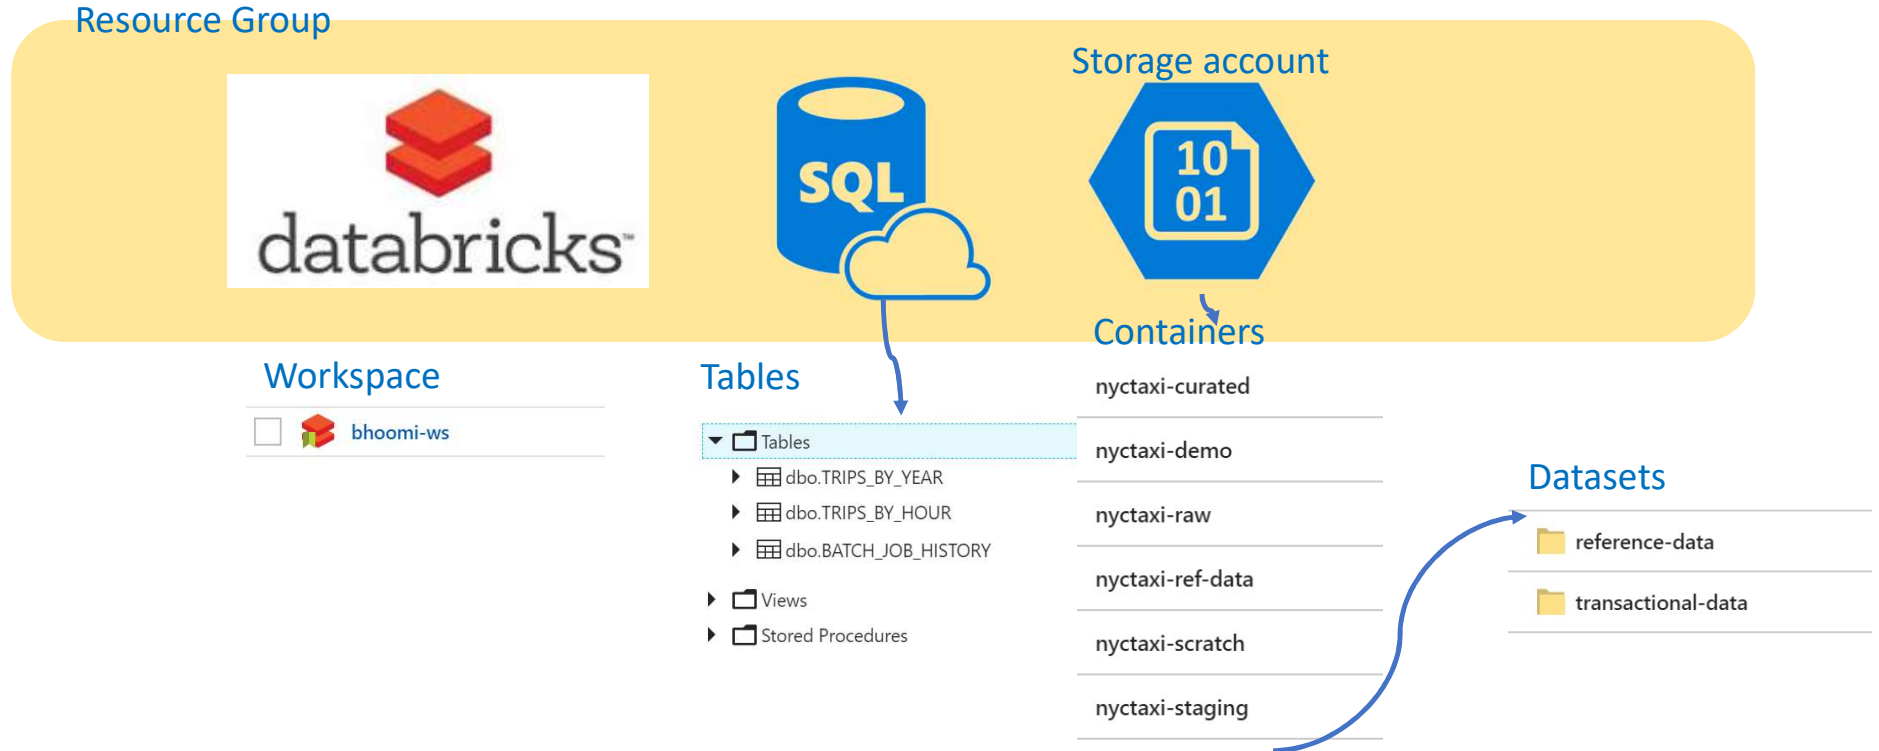


**Instructions:**

<https://aka.ms/DEW-Setup>

# By now...

You should have the following set up in your subscription:



# **Azure Databricks**

## **Provisioning, configuration**

# Provision cluster

## Azure Databricks – setup

Initialize workspace and create a cluster with the following specifications:

Cluster Type

☐ Serverless Pool (beta, R/Python/SQL) ☒ Standard [Learn more about Serverless Pools ?](#)

Databricks Runtime Version

Python Version ?

Driver Type

56.0 GB Memory, 8 Cores, 2 DBU

Worker Type

56.0 GB Memory, 8 Cores, 2 DBU

Min Workers  Max Workers  ☒ Enable Autoscaling ?

Auto Termination ?

☒ Terminate after  minutes of inactivity

In the Spark config section, enter your storage account credentials-  
spark.hadoop.fs.azure.account.key.<storageAccount>.blob.core.windows.net <key>

# By now...

You should have the following set up in your subscription:

## Resource Group



## Storage account



## Containers

## Workspace & cluster

☐ bhoomi-ws

Clusters

[+ Create Cluster](#)

▼ Interactive Clusters

Name	State	Nodes
bhoomi	Running	4

## Tables

- ▼ Tables
  - ▶ dbo.TRIPS\_BY\_YEAR
  - ▶ dbo.TRIPS\_BY\_HOUR
  - ▶ dbo.BATCH\_JOB\_HISTORY
- ▶ Views
- ▶ Stored Procedures

nyctaxi-curved

nyctaxi-demo

nyctaxi-raw

nyctaxi-ref-data

nyctaxi-scratch

nyctaxi-staging

## Datasets

reference-data

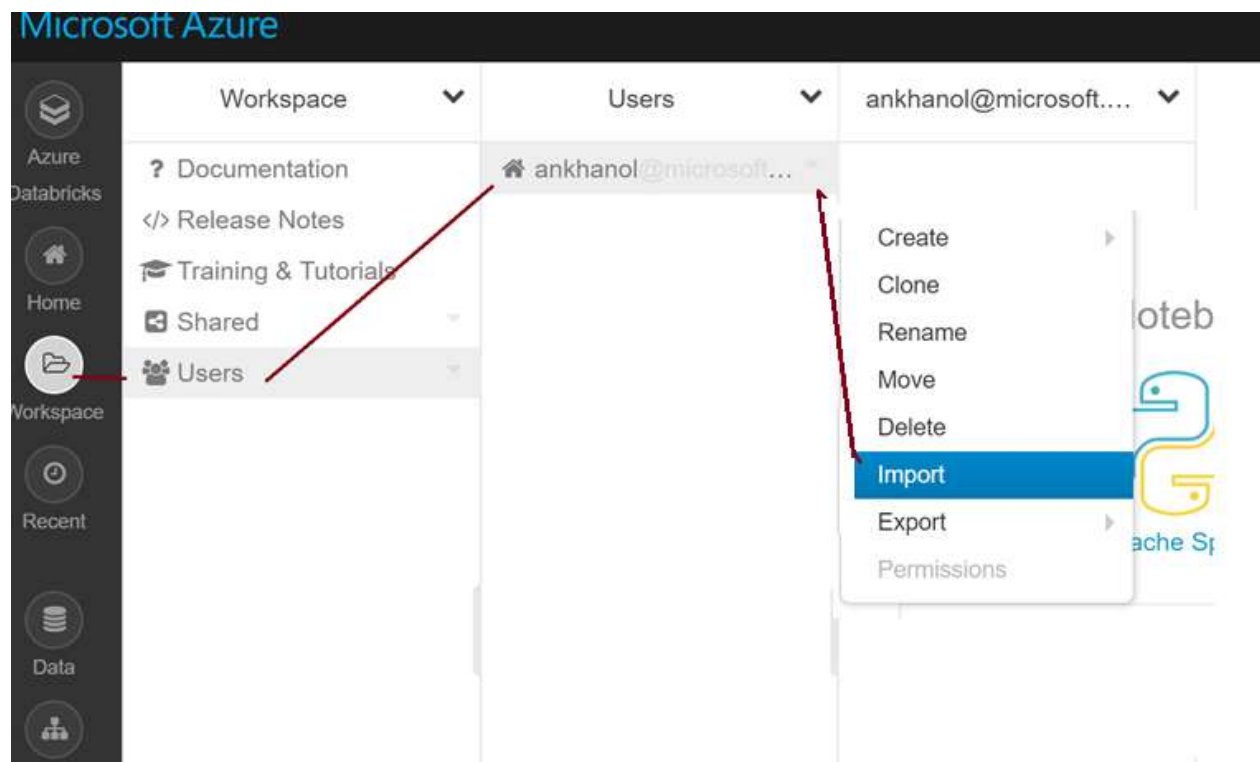
transactional-data

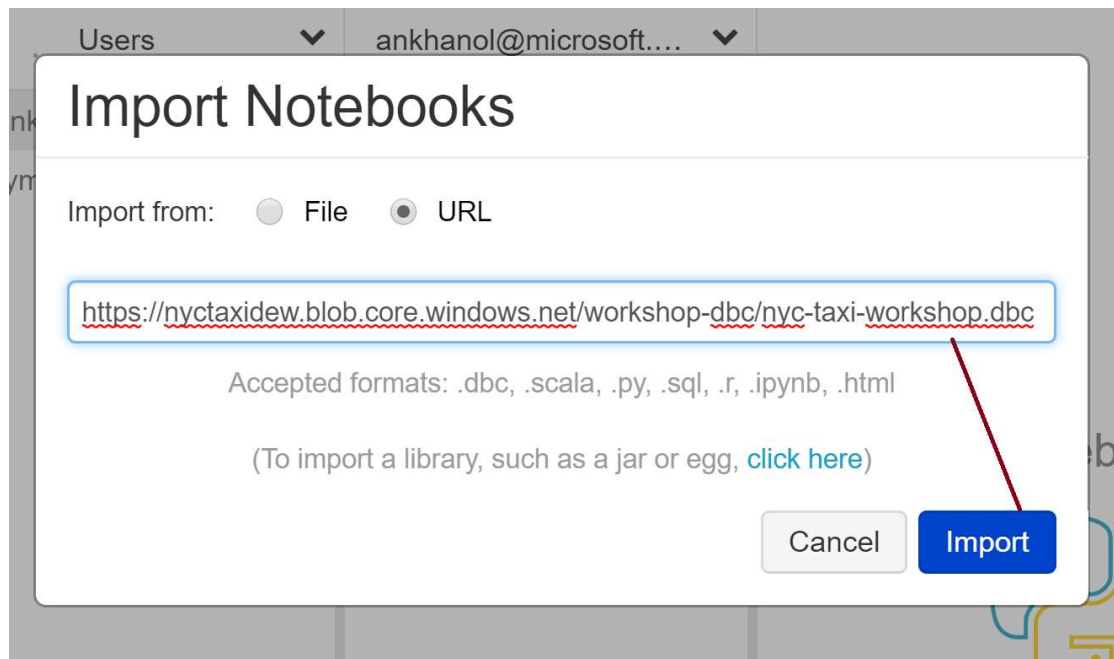
# Azure Databricks

## Import workshop notebooks





















# Import workshop notebooks





<https://nyctaxidew.blob.core.windows.net/workshop-dbc/nyc-taxi-workshop.dbc>

<div><div> Azure</div><div> Databricks</div><div> Home</div><div> Workspace</div><div> Recent</div></div>	Workspace ▼	Users ▼	ankhanol@microsoft.... ▼	nyc-taxi-workshop ▼	00-HowTo ▼
	? Documentation	 ankhanol@microsoft...		 00-HowTo ▼	 1-MountBlobStorage ▼
	</> Release Notes			 01-General ▼	 2-DownloadDataFro... ▼
	🎓 Training & Tutorials			 02-LoadData ▼	 3-UseSparkSQLWith... ▼
	📁 Shared ▼		 nyc-taxi-workshop ▼	 03-TransformData ▼	 4-WorkingWithRemo... ▼
	 Users ▼			 04-CreateMaterializ... ▼	 5-WorkingWithStorage ▼

# You are done with setup!

You should have the following set up in your subscription:

## Resource Group



## Storage account



## Containers

nyctaxi-curved

nyctaxi-demo

nyctaxi-raw

nyctaxi-ref-data

nyctaxi-scratch

nyctaxi-staging

## Datasets

reference-data

transactional-data

## Workspace, cluster, notebooks

nyctaxi-workshop

- 00-HowTo
- 01-General
- 02-LoadData
- 03-TransformData
- 04-CreateMaterializ...
- 05-GenerateReports
- 06-BatchJob

Clusters

+ Create Cluster

Interactive Clusters

Name	State	Nodes
bhoomi	Running	4

## Tables

- Tables
  - dbo.TRIPS\_BY\_YEAR
  - dbo.TRIPS\_BY\_HOUR
  - dbo.BATCH\_JOB\_HISTORY
- Views
- Stored Procedures

# Password for demo database

The password for the RDBMS demodbserver is:  
d@t@br1ck\$

# **Azure Databricks**


## **Preview of what's you'll be creating**


# Query-able Hive tables


	Databases ▼	Tables +
	🔍 Filter Databases	🔍 Filter Tables
	🗄 default	🗄 green_taxi_trips ▼
	🗄 demo_db	🗄 green_taxi_trips_cur... ▼
	🗄 taxi_db	🗄 payment_type_lookup ▼
	🗄 taxi_reports_db	🗄 rate_code_lookup ▼
		🗄 taxi_trips_mat_view ▼
		🗄 taxi_zone_lookup ▼
		🗄 trip_month_lookup ▼
		🗄 trip_type_lookup ▼
		🗄 vendor_lookup ▼
		🗄 yellow_taxi_trips ▼
		🗄 yellow_taxi_trips_cu... ▼


# On-demand/scheduled batch jobset


Microsoft Azure


 Azure


 Databricks


 Home


 Workspace

 Recent

 Data


 Clusters

 Jobs

 Search

NYCTaxiReportJobSet

[< All Jobs](#)



✕ Delete

Job ID: 14

Task: Notebook at [/Users/ankhanol@microsoft.com/nyc-taxi/06-BatchJob/Workflow](#) - [Edit](#) / [Remove](#)

▸ Parameters: [Edit](#)

◦ Dependent Libraries: [Add](#)

Cluster: bhoomi (224 GB, Running, 4.0 beta (Scala 2.11)) [Edit](#)

Schedule: None [Edit](#)

Advanced ▸

Active runs

Run	Start Time	Launched	Duration
<a href="#">Run Now</a> / <a href="#">Run Now With Different Parameters</a>			

Completed in past 60 days

[Latest successful run \(refreshes automatically\)](#)

< Previous 20

Run	Start Time	Launched	Duration
<a href="#">Run 10</a>	2018-02-19 16:26:55 CST	Manually	54s
<a href="#">Run 9</a>	2018-02-19 16:25:38 CST	Manually	54s



# Hands on lab - Module 1

## HOW TO

work with Hive, with DBFS, with remote databases

# 1.1. Mount blob storage

Why?

What's easier?

Option 1: wasbs URI

`wasbs://storageContainer@storageAccount.blob.core.windows.net/<myDirectory>`

Option 2: Mount point

`/mnt/data/<myDirectory>`

Mounting blob storage is an option with Azure Databricks 3.5 and above, and a general best practice.

It simplifies and secures accessing storage

# 1.1. Mount blob storage

How to

Lets review the notebook –  
`nyc-taxi-workshop/00-HowTo/00-MountBlobStorage.scala`

In this notebook – we will mount a storage container, and access it using the mountpoint. We will also learn to refresh mountpoints and to unmount.

## 1.2. Downloading the original Yellow Taxi data

How to

**This is informational only**

Lets review the notebook –  
`nyc-taxi-workshop/00-HowTo/2-DownloadDataFromInternet.scala`

## 1.3. Working with Hive & SparkSQL

How to

Lets review the notebook –

`nyc-taxi-workshop/00-HowTo/3-UseSparkSQLWithHive.scala`

In this notebook – we will create a text file locally (bash), load it to DBFS, create a hive table on it, and run queries using SQL

## 1.4. Working with remote databases/JDBC

How to

Lets review the notebook –

`nyc-taxi-workshop/00-HowTo/4-WorkingWithRemoteDatabases.scala`

In this notebook – we will query a remote Azure SQL database like it were a table created in Azure Databricks

## 1.5. Working with storage

How to

Lets review the notebook –

`nyc-taxi-workshop/00-HowTo/5-WorkingWithStorage.scala`

In this notebook, we will work with DBFS file system commands

# Recap of module

We learned how to -

- 1) Mount blob storage
- 2) Work with Hive and SparkSQL
- 3) Work with a remote SQL database
- 4) Work with Databricks File System (DBFS)



# Hands on lab - Module 2

General setup

Mount blob storage

Create database

Common functions

## 2.1. Mount blob storage

### Workshop - setup

Lets review this notebook-  
`nyc-taxi-workshop/01-General/1-MountBlobStorage.scala`

In this notebook, we create a function to mount blob storage and call it to mount several containers

## 2.2. Create database objects

### Workshop - setup

Lets review this notebook-  
`nyc-taxi-workshop/01-General/2-CreateDatabaseObjects.scala`

In this notebook, -

- (1) We create a database `taxi_db`
- (2) We create a JDBC Hive table against the Azure SQL database table you set up during provisioning – `batch_job_history`

## 2.3. Define common functions

Focus - reusability

Lets review this notebook-  
`nyc-taxi-workshop/01-General/3-CommonFunctions.scala`

In this notebook, -  
We create a set of commonly used functions in a notebook for future use in the workshop

# Recap of module

We completed-

- 1) Mounting blob storage
- 2) Creating a Hive database
- 3) Creating a Hive table definition for a remote Azure SQL database
- 4) Creating a set of commonly used functions in a notebook for future use in the workshop

# Hands on lab - Module 3


Load dataset

## 3.1. Load reference data

Lets review this notebook-  
`nyc-taxi-workshop/02-LoadData/1-LoadReferenceData.scala`

In this notebook, -

- (1) We load all the reference datasets – read from stagingDir, save as parquet to refDataDir
- (2) Create Hive tables on top of them
- (3) Compute statistics



Trip months  
Payment type  
Rate code  
Taxi zone  
Trip type  
Vendor

## 3.2. Load transactional data – Yellow Taxi

Lets review this notebook-

`nyc-taxi-workshop/02-LoadData/2-LoadData-YellowTaxi.scala`

In this notebook, -

- (1) We load yellow taxi data over several years– read from stagingDir, save as parquet to refDataDir
- (2) Create Hive tables on top of them
- (3) Compute statistics

To note: the schema of the dataset differs over years. We homogenize the schema in this notebook



## 3.3. Load transactional data – Green Taxi

Lets review this notebook-  
`nyc-taxi-workshop/02-LoadData/3-LoadData-GreenTaxi.scala`

In this notebook, -

- (1) We load green taxi data over several years– read from stagingDir, save as parquet to refDataDir
- (2) Create Hive tables on top of them
- (3) Compute statistics

To note: the schema of the dataset differs over years. We homogenize the schema in this notebook

## Recap of module

We completed-

- 1) Loading reference data and saving as parquet, with hive tables & stats
- 2) Loading yellow taxi and green taxi data, saving as parquet, with hive tables & stats

# Hands on lab - Module 4

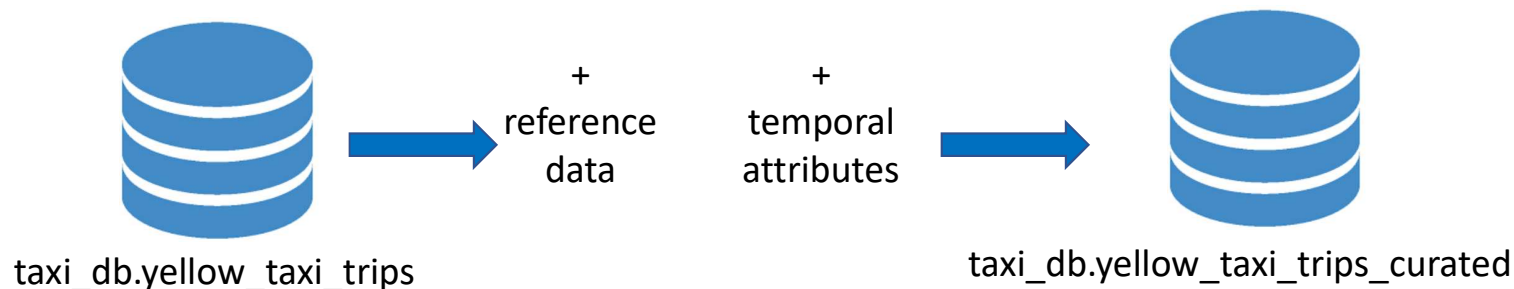
Transform data

## 4.1. Transform yellow taxi data

Lets review this notebook-  
`nyc-taxi-workshop/03-TransformData/1-TransformData-YellowTaxi.scala`

In this notebook, -

- (1) We execute the notebook with common functions
- (2) Denormalize - join yellow taxi with all reference datasets and persist to DBFS
- (3) Define Hive table for the transformed dataset (curated), create partitions (year, month)
- (4) Compute statistics




## 4.2. Transform green taxi data

Lets review this notebook-

`nyc-taxi-workshop/03-TransformData/2-TransformData-GreenTaxi.scala`

In this notebook, -

- (1) We execute the notebook with common functions
- (2) Denormalize - join yellow taxi with all reference datasets and persist to DBFS
- (3) Define Hive table for the transformed dataset (curated), create partitions (year, month)
- (4) Compute statistics



## Recap of module

We completed-

Denormalizing yellow and green trip taxi data – joining with reference data, and saving as parquet, and ran hive tables & stats

# Hands on lab - Module 5

Create materialized view

## 5.1. Create materialized view

Lets review this notebook-  
[nyc-taxi-workshop/04-CreateMaterializedView/1-CreateMaterializedViews.scala](#)

In this notebook, -

- (1) We execute the notebook with common functions
- (2) We add columns to yellow taxi data to match green taxi data
- (3) Do a union all on yellow and green taxi curated datasets and persist to DBFS
- (4) Define Hive table for the transformed dataset (curated), create partitions (taxi type, year, month)
- (5) Compute statistics





# Hands on lab - Module 6

Generate a report with visualization

## 6.1. Generate report

Lets review this notebook-  
`nyc-taxi-workshop/05-GenerateReports/Report-1.scala`

In this notebook, we will run several reports and visualize-

1. Trip count by taxi type
2. Revenue including tips by taxi type
3. Revenue share by taxi type
4. Trip count trend between 2013 and 2016
5. Trip count trend by month, by taxi type, for 2016
6. Average trip distance by taxi type

## 6.1. Generate report

7. Average trip amount by taxi type
8. Trips with no tip, by taxi type
9. Trips with no charge, by taxi type
10. Trips by payment type
11. Trip trend by pickup hour for yellow taxi in 2016
12. Top 3 yellow taxi pickup-dropoff zones for 2016

# Hands on lab - Module 7

Batch job automation

## 7.1. Global vars and functions

Lets review the notebook-  
`nyc-taxi-workshop/06-BatchJob/GlobalVarsAndMethods.scala`

In this notebook-

- (1) We define JDBC credentials
- (2) We create a function to generate a batch\_ID by querying the batch\_job\_history table in Azure SQL database

## 7.2. Report 1

Lets review the notebook-  
`nyc-taxi-workshop/06-BatchJob/Report-1.scala`

In this notebook-

- (1) We create a dataframe – a simple report – trips by year
- (2) Persist the data to Azure SQL database table `trips_by_year`

## 7.3. Report 2

Lets review the notebook-  
`nyc-taxi-workshop/06-BatchJob/Report-2.scala`

In this notebook-

- (1) We create a dataframe – a simple report – trips by hour
- (2) Persist the data to Azure SQL database table `trips_by_hour`

## 7.4. Workflow

Lets review the notebook-  
nyc-taxi-workshop/06-BatchJob/Workflow.scala


This is a notebook workflow spec and in this we –

- (1) Execute the GlobalVarsAndFunctions notebook
- (2) We generate batch ID by calling a method from #1
- (3) We insert start time into batch\_job\_history, execute notebook Report-1, then insert completion time into the RDBMS table




## 7.4. Workflow


- (4) If #3 completed successfully, we repeat #3 for Report-2
- (5) If #4 completed successfully, we exit with a pass status
- (6) We create a batch job in which we call `workflow.scala` and do multiple ad-hoc executions




Azure  
Databricks



Home



Workspace



Jobs

+ Create Job

Name ↑	Job ID	Created By	Task
● <a href="#">FlightDelayPrediction</a>	5	Anagha Khanolkar	<a href="#">8.3-Workflow</a>
● <a href="#">NYCTaxiReportJobSet</a>	14	Anagha Khanolkar	<a href="#">Workflow</a>

Azure  
Databricks

Home

Workspace

Recent

Data

Clusters

Jobs

Search

## NYCTaxiReportJobSet

[< All Jobs](#)

## NYCTaxiReportJobSet



✕ Delete

Job ID: 14

Task: Notebook at </Users/ankhanol@microsoft.com/nyc-taxi/06-BatchJob/Workflow> - [Edit](#) / [Remove](#)

▸ Parameters: [Edit](#)

◦ Dependent Libraries: [Add](#)

Cluster: bhoomi (224 GB, Running, 4.0 beta (Scala 2.11)) [Edit](#)

Schedule: None [Edit](#)

Advanced ▸

## Active runs

Run	Start Time	Launched	Duration
<a href="#">Run Now</a> / <a href="#">Run Now With Different Parameters</a>			

## Completed in past 60 days

[Latest successful run \(refreshes automatically\)](#)

< Previous 20

Run	Start Time	Launched	Duration
<a href="#">Run 10</a>	2018-02-19 16:26:55 CST	Manually	54s
<a href="#">Run 9</a>	2018-02-19 16:25:38 CST	Manually	54s
<a href="#">Run 8</a>	2018-02-19 16:13:05 CST	Manually	52s

**Wrap up**

# Machine Learning Workshop - preview

## Lecture/discussion:

Data science process  
Data science lexicon – level set

## Lab:

Use case: Flight delay prediction

## Technologies/Services:

(1) Azure Machine Learning (AML) Studio – we will learn how to predict flight delays using AML – ingest, cleanse, dedupe, train, test, operational model, batch score

(2) Spark ML – we will repeat the exact same experiment in Spark ML – dataframe API for a more scalable solution

We will discuss how we can operationalize a Spark model trained in Databricks on the Azure ML platform

## Q & A

1. Terminate your cluster
2. Any questions?

Note: you will use the same cluster next week for the machine learning on Databricks workshop